

# Exercise7

January 10, 2019

Machine Learning Lab - Exercise Sheet 7

Author: Sebastian Pineda Arango ID: 246098 Universität Hildesheim - Data Analytics Master

## 0.1 Exercise sheet 7

In this exercise sheet, two methods are examined: K-Nearest neighbours and decision trees. The data set used to train the models is the classical iris data set [1].

The data set is imported from the module *datasets* include in scikit-learn. We also explore the shape, some descriptive statistics and the distributions of the labels of the dataset. Afterwards, some insights are extracted from the dataset:

- There are 150 samples
- 4 features (corresponding to width and length of petals and sepals.
- The classes are balanced: the number of samples for each class is roughly the same
- The descriptors change the scale, therefore it is necessary to rescale so that distance based methods perform better.

```
In [19]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
import statistics as st
from sklearn.base import BaseEstimator
```

```
%matplotlib inline
```


```
In [20]: data = datasets.load_iris()
X = data.data
X.shape
```

```
Out[20]: (150, 4)
```

```
In [21]: Y = data.target
Y.shape
```

```
Out[21]: (150,)
```

```
In [22]: plt.hist(Y)
plt.grid()
```



```
In [23]: pd.DataFrame(X).describe()
```

```
Out [23]:
```

	0	1	2	3
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

### 0.1.1 Exercise 1: Implement K-Nearest Neighbor (KNN)

Before proceeding, two operations are applied over the dataset:

- We divide data set in training and test set (70/30).
- We fit a standard scaler over the training set and we transform it on test set. To check that it was rescaled, we get the basic statistics again. Afterwards, the mean is effectively zero and the standard deviation is one.

```
In [24]: from sklearn.model_selection import train_test_split
         from sklearn.preprocessing import StandardScaler

         #Splitting train-test set
         X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3)
```

```

scaler = StandardScaler()

#Standardizing data
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

#Finding main descriptors of the data
pd.DataFrame(X_train_scaled).describe()

```

```

Out [24]:

```

	0	1	2	3
count	1.050000e+02	1.050000e+02	1.050000e+02	1.050000e+02
mean	-5.202188e-16	6.952111e-16	2.008975e-16	-5.075305e-16
std	1.004796e+00	1.004796e+00	1.004796e+00	1.004796e+00
min	-1.796993e+00	-2.409310e+00	-1.454129e+00	-1.382159e+00
25%	-8.553371e-01	-5.844044e-01	-1.231070e+00	-1.254406e+00
50%	-3.138852e-02	-1.281779e-01	3.303390e-01	1.508694e-01
75%	6.748531e-01	5.561618e-01	7.764560e-01	7.896312e-01
max	2.440457e+00	3.065407e+00	1.780219e+00	1.683898e+00

A class is implemented that wraps many functionalities of the KNN classifier. This class inherits from Scikit-learn BaseEstimator (so that it can be used along with other scikit-learn functionalities). The class defines also several methods:

- Constructor (**init** method): used to initialize k when the object is created.
- fit: to train the model. Since KNN has no model, we storage the training and test set.
- predict: to make predictions over a given test set.
- score: to find the accuracy over a test set.
- find knn: return a list of nearest neighbours
- euclidean distance: find the euclidean distance between different rows of a matrix and a query vector.

Accuracy is chosen as performance measure, since the label distribution is balanced. Otherwise, we should use another measure, less sensible to label distribution (lik F1-Score). We can define accuracy as:

$$Accuracy = \frac{TP+TN}{S}$$

Where TP are True Positives, TN are True Negatives and S is the size of the whole set. We could also define accuracy as:

$$Accuracy = \sum I(y = \hat{y})$$

Where  $I = 1$  if  $y = \hat{y}$ , and  $I = 0$  otherwise.

Finally, an object is created to make predictions over a test set using KNN (k=3) and then, we find, the accuracy over a test set.

```

In [25]: #class knn estimator

```

```

class knn_estimator (BaseEstimator):

    def __init__(self,k):

        '''Intialize the hyperparameter k (number of nearest neighbors to look for)'''

```

```

        self.k=k

def fit(self, X_train, y_train):

    '''Save the train set (features and labels) to be able to make predictions in

    self.X_train = X_train
    self.y_train = y_train

def predict(self, X_test):

    '''Make predictions over the X_test using KNN'''

    n_test_samples = X_test.shape[0]
    predictions = []

    for i in range(n_test_samples):
        X_query = X_test[i,:]
        (top_k, idx) = self.find_knn(self.X_train, X_query, self.k)
        #pred_class = st.mode(self.y_train[idx,])
        pred_class= max(set(list(self.y_train[idx,].tolist())), key=self.y_train.t
        predictions.append(pred_class)

    return np.array(predictions)

def score(self, X, y):

    '''Find the accuracy oF the predictions over X given the y labels'''

    y_pred = self.predict(X)
    acc=np.mean(y==y_pred)

    return acc

def find_knn(self, X, X_query, k):

    '''Find the k-nearest neighbors of X_query among the samples in X '''

    dist = self.euclidean_distance(X, X_query)
    idx = np.argsort(dist.T[0], axis=0)[:k]
    top_k = X[idx,:]
    return top_k, idx

def euclidean_distance(self, X, X_query):

    '''Find the euclidean distance between X_query (single sample) and X
    (matrix where each row is a sample)'''

```

```

if X.shape[1] != X_query.shape[0]:
    print("Dimensions doesn't match.")

dist = np.zeros((X.shape[0], 1))

for i in range(X.shape[0]):
    dist[i] = np.sqrt(np.sum((X[i,:]-X_query)**2))

return dist

```

```

k=5
knn_clf = knn_estimator(k)
knn_clf.fit(X_train, y_train)
acc = knn_clf.score(X_test, y_test)
acc

```

Out [25]: 0.9333333333333333

### 0.1.2 Exercise 2 - Part A: Optimize and compare KNN algorithm

To determine the optimal value of K in KNN, we will use grid search with cross validation. For that we apply k-fold cross validation [2] and tried a different set of hyperparameters. At the end, we choose the hyperparameter with the highest mean accuracy for the k-fold.

Note that the hyperparameter that we aim to fin is the number of neighbors (K).

```

In [26]: #number of samples of training set
n_train = X_train.shape[0]

#number of folds
n_folds = 5

#initializing folds
folds = []
samples_fold = int(n_train/n_folds)

#creating the k-fold subsets
for i in range(n_folds):

    folds.append((X_train[(i*samples_fold):((i+1)*samples_fold)],
                    y_train[(i*samples_fold):((i+1)*samples_fold)]))

folds_list = list(range(n_folds+1))

#initialize list to store the man of each hyperparameter setting
mean_test_folds = []

```

```

#list of hyperparameters
k_list = [1,2,3,5,10,20]

for k in k_list:

    test_acc_folds = []

    for f in range(n_folds):

        #list of folds
        folds_list = list(range(n_folds))
        folds_list.pop(f)

        #selecting test dataset
        X_test_fold = folds[f][0]
        y_test_fold = folds[f][1]

        #merging the folds to create the training dataset
        X_train_fold = folds[folds_list[1]][0]
        y_train_fold = folds[folds_list[1]][1]

        for j in folds_list[1:]:
            X_train_fold = np.vstack((X_train_fold, folds[j][0]))
            y_train_fold = np.hstack((y_train_fold, folds[j][1]))

        #creating base estimator
        knn_clf = knn_estimator(k)

        #saving the X_train and y_train
        knn_clf.fit(X_train_fold, y_train_fold)


        #finding accuracy over the fold
        acc_test = knn_clf.score(X_test_fold, y_test_fold)
        test_acc_folds.append(acc_test)

    #findning the mean across all the folds
    mean_test_folds.append(np.mean(test_acc_folds))

plt.plot(k_list, mean_test_folds, '*')
plt.grid()
plt.xlabel("k")
plt.ylabel("Accuracy")

```

Out[26]: Text(0,0.5,'Accuracy')



After running cross validation, we can see that the best value for the hyperparameters is  $K=1$ . We also validate this result using *GridSearch* function of Scikit learn. We can do this since we have created the *knn\_estimator* class with the same interface as an estimator class of Scikit-learn. This also returns that the best value for  $K$  is 1.

```
In [27]: from sklearn.model_selection import GridSearchCV
```

```
knn_clf = knn_estimator(3)
parameters={'k':k_list}
knn_grid = GridSearchCV(knn_clf, parameters, cv=5)
knn_grid.fit(X_train, y_train)
best_knn= knn_grid.best_estimator_
best_knn
```

```
Out[27]: knn_estimator(k=2)
```

To give evidence that this value is the better, we try on test set different values for  $K$ . As we can see, the greater the value of  $K$ , the lower the accuracy. The graph show that the best and lowest value for  $K$  is  $K=1$ .


```
In [28]: acc_list = []
```

```
for i in range(1,25):
    knn_clf = knn_estimator(i)
    knn_clf.fit(X_train, y_train)
    acc = knn_clf.score(X_test, y_test)
```

```

    acc_list.append(acc)
plt.plot(acc_list)
plt.grid()

```



### 0.1.3 Exercise 2 - Part B: Comparison of KNN with Tree based methods

The performance of decision tree and KNN is compared using the test set and analyzing the stability of the results using `cross_val_score`. The idea is that the results do not be very deviated from a mean value.

Hyperparameter search is applied over the decision trees. In this case, we vary maximal depth of the tree as hyperparameter. We can see that the best tree is the one with `max_depth = 3`.

```

In [31]: from sklearn import tree

parameters = {'max_depth':list(range(1,11))}

#creating the classifier
tree_clf = tree.DecisionTreeClassifier()
tree_grid = GridSearchCV(tree_clf, parameters, cv=5)

#performing hyperparameter search
tree_grid.fit(X_train, y_train)

#finding the best decision tree
best_tree = tree_grid.best_estimator_

```



```

print("Best tree:", best_tree)

#plotting differen results
x_axis = [i['max_depth'] for i in tree_grid.cv_results_['params']]
x_axis
y_axis = tree_grid.cv_results_['mean_test_score']

plt.plot(x_axis, y_axis)
plt.grid()
plt.xlabel("Max. depth")
plt.ylabel("Accuracy")

```

```


Best tree: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=3,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, presort=False, random_state=None,
    splitter='best')

```

```

Out[31]: Text(0,0.5,'Accuracy')

```



```

In [32]: from sklearn.model_selection import cross_val_score

metrics_tree = cross_val_score(best_tree, X_train, y_train, cv=5)

```


```
plt.boxplot(metrics_tree)
plt.grid()

print("Results for the best tree:")
print("Mean of results:", np.mean(metrics_tree))
print("Standard deviation of results:", np.std(metrics_tree))
```

Results for the best tree:

Mean of results: 0.95

Standard deviation of results: 0.04472135955



```
In [33]: from sklearn.model_selection import cross_val_score


metrics_knn = cross_val_score(best_knn, X_train, y_train, cv=5)
plt.boxplot(metrics_knn)
plt.grid()

print("Results for the best KNN classifier:")
print("Mean of results:", np.mean(metrics_knn))
print("Standard deviation of results:", np.std(metrics_knn))
```

Results for the best KNN classifier:

Mean of results: 0.952380952381

Standard deviation of results: 0.042591771



```
In [34]: #Evaluation on test set
acc_knn_test = best_knn.score(X_test, y_test)
acc_tree_test = best_tree.score(X_test, y_test)

print("Accuracy of KNN:", acc_knn_test)
print("Accuracy of tree:", acc_tree_test)
```

```
Accuracy of KNN: 0.955555555556
Accuracy of tree: 0.933333333333
```

The KNN classifier performs *slightly* better than decision tree. In fact, during the creation of this notebook, the results have changed a little bit between different runs. In this case, we can see that the accuracy on test set is better for KNN, however, due to the small size of test set (45 samples), this difference only means one wrongly classified sample.

At the end of the day, although the KNN seems to perform *slightly* better (at least for this run), the decision tree could be better in other aspects. For instance, decision tree stores few parameters to make predictions, while KNN has to store the whole training set. Moreover, in spite of a lacking execution time comparisons, it is highly probable that the decision tree is faster than KNN, since the former has only to evaluate a few conditions, while the last one has to calculate the euclidean distance respect to every training sample.

## 0.2 References

- Iris data set: [https://en.wikipedia.org/wiki/Iris\\_flower\\_data\\_set](https://en.wikipedia.org/wiki/Iris_flower_data_set)
- K-Fold cross validation: <https://machinelearningmastery.com/k-fold-cross-validation/>