

# Exercise7\_Solution

June 23, 2019

## 1 Distributed Data Analytics

## 2 Exercise Sheet 6

### 2.1 Sebastian Pineda Arango Mtr. Nr. 246098

In this notebooks we are going to implement two Convolutional Neural Networks, a complex and simple one. Moreover, we are going to apply data transformations to increase the size of the training data. The models will be trained on CIFAR10 dataset [3].

#### 2.1.1 Data Augmentation

For the data augmentation process, we are going to implement the following transformations using Tensorflow:

- Scaling
- Translation
- Rotation

To code the transformations, we have partially taken references from [4] and [5].

```
In [1]: import numpy as np
import tensorflow as tf
import keras
import math
from keras.datasets import cifar10
import matplotlib.pyplot as plt
import os

%matplotlib inline
```

```
C:\Users\User\Anaconda3\lib\site-packages\h5py\__init__.py:34: FutureWarning: Conversion of the
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [2]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```


num_images = x_train.shape[0]
image_size = x_train.shape[1]
num_classes = np.max(y_train)+1

```

```

#loading image sample
plt.figure()
plt.imshow(x_train[1,:,:,:])
plt.show()
print(x_train.shape)

```



(50000, 32, 32, 3)

```

In [3]: X = tf.placeholder(tf.uint8, shape = (None, image_size, image_size, 3))
        y = tf.placeholder(tf.uint8, shape=(None, 1), name='output_y')

```

```

class DataAugmentation():

    def __init__(self, image_size, num_images):

        """Contructor for the data augmentation class"""

        self.image_size = image_size
        self.num_images = num_images

        print("Creating Data Agumentation Class")

```

```

def rotate(self, angle=30):

    """Rotation of the image"""

    rotate_tf = tf.dtypes.cast(tf.contrib.image.rotate(self.X, angle), tf.uint8)
    self.rotate_tf = rotate_tf


def scale(self, scale=0.8):

    """Scaling of the image"""

    n_imgs = self.num_images
    boxes = np.zeros((n_imgs, 4), dtype = np.float32)

    for index in range(n_imgs):
        x1 = y1 = 0.5 - 0.5 * scale # To scale centrally
        x2 = y2 = 0.5 + 0.5 * scale
        boxes[index] = np.array([y1, x1, y2, x2], dtype = np.float32)
    box_ind = np.arange(n_imgs)
    crop_size = np.array([image_size, image_size], dtype = np.int32)
    scale_tf = tf.dtypes.cast(tf.image.crop_and_resize(self.X, boxes, box_ind, crop_size), tf.uint8)
    self.scale_tf = scale_tf


def translate(self, pix_trans=10):

    """Translation of the image"""

    translate_tf = tf.dtypes.cast(tf.contrib.image.translate(self.X, translations=[pix_trans, pix_trans]), tf.uint8)
    self.translate_tf = translate_tf


def get_rotate(self):

    return self.rotate_tf


def get_scale(self):

    return self.scale_tf


def get_translate(self):

    return self.translate_tf


def augment_data(self, X, y, angle, scale, pix_trans):

    """Main function to augment the data"""

```

```

        angle_radians = np.pi*(angle/180) #transforming to radians
        self.X = X
        self.rotate(angle_radians)
        self.scale(scale)
        self.translate(pix_trans)

        return tf.concat([self.X, self.rotate_tf, self.scale_tf, self.translate_tf], a

scale=0.9
pix_trans = 10
angle = 20
data_augmenter = DataAugmentation(image_size=image_size, num_images=num_images)
da_tensor, da_labels = data_augmenter.augment_data(X, y, angle, scale, pix_trans)

sess = tf.Session()

X_train_a, y_train_a, ex_rotate, ex_scale, ex_translate = sess.run([da_tensor, da_labels,
                                                                    data_augmenter.get_scale()[0], data_augmenter.get_translate()[0],
                                                                    feed_dict= {X:x_train,
                                                                    y: y_train})

sess.close()

```

Creating Data Augmentation Class

WARNING: The TensorFlow contrib module will not be included in TensorFlow 2.0.

For more information, please see:

- \* <https://github.com/tensorflow/community/blob/master/rfcs/20180907-contrib-sunset.md>
- \* <https://github.com/tensorflow/addons>

If you depend on functionality not listed there, please file an issue.

In [4]: *#checking the size after data augmentation*

```

print(X_train_a.shape)
print(y_train_a.shape)

```

```

(200000, 32, 32, 3)

```

```

(200000, 1)

```

In [5]: *#checking data augmentation*

```

plt.imshow(x_train[0,:,:,:])
print("Original")
plt.show()


plt.imshow(ex_rotate)
print("Example from rotation")
plt.show()

```


```
plt.imshow(ex_scale)
print("Example from scaling")
plt.show()

plt.imshow(ex_translate)
print("Example from translation")
plt.show()
```


Original




Example from rotation



Example from scaling



Example from translation



### 2.1.2 Simple convolutional network

We implement a simple neural network with only a convolutional layer, a max-pooling layer and a fully connected layer. For *regularizing* the neural network, there are several techniques:

- Data augmentation: it is, generating new images by transforming the original training images (as we have done before)
- Drop-out: by eliminating, randomly, some activations.
- Batch-normalization: which has been reported to have a normalizing effect [6].
- L2 (or L1) normalization: by penalizing the weights magnitude in the loss function.

In the following examples we are going to apply L2 (in the simple model) and Batch Normalization (in the complex model).

```
In [8]: def random_mini_batches(X, Y, mini_batch_size = 64, seed = 0):
        """
        Creates a list of random minibatches from (X, Y)

        Arguments:
        X -- input data, of shape (input size, number of examples)
        Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of shape (1, number of examples)
        mini_batch_size -- size of the mini-batches, integer
        seed -- this is only for the purpose of grading, so that you're "random minibatches"
```

*Returns:*

*mini\_batches -- list of synchronous (mini\_batch\_X, mini\_batch\_Y)*  
"""

```
m = X.shape[0]                # number of training examples
mini_batches = []
np.random.seed(seed)
```

*# Step 1: Shuffle (X, Y)*

```
permutation = list(np.random.permutation(m))
shuffled_X = X[permutation,:]
shuffled_Y = Y[permutation,:]
```

*# Step 2: Partition (shuffled\_X, shuffled\_Y). Minus the end case.*

```
num_complete_minibatches = math.floor(m/mini_batch_size)
```

*# number of mini batches of size mini\_batch\_size in your partitionning*

```
for k in range(0, num_complete_minibatches):
```

```
    mini_batch_X = shuffled_X[k * mini_batch_size : k * mini_batch_size + mini_batch_size]
```

```
    mini_batch_Y = shuffled_Y[k * mini_batch_size : k * mini_batch_size + mini_batch_size]
```

```
    mini_batch = (mini_batch_X, mini_batch_Y)
```

```
    mini_batches.append(mini_batch)
```

*# Handling the end case (last mini-batch < mini\_batch\_size)*

```
if m % mini_batch_size != 0:
```

```
    mini_batch_X = shuffled_X[num_complete_minibatches * mini_batch_size : m,:]
```

```
    mini_batch_Y = shuffled_Y[num_complete_minibatches * mini_batch_size : m, :]
```

```
    mini_batch = (mini_batch_X, mini_batch_Y)
```

```
    mini_batches.append(mini_batch)
```

```
return mini_batches
```

```
def variable_summaries(var):
```

```
    """Attach a lot of summaries to a Tensor (for TensorBoard visualization)."""
```

```
    with tf.name_scope('summaries'):
```

```
        mean = tf.reduce_mean(var)
```

```
        tf.summary.scalar('mean', mean)
```

```
        with tf.name_scope('stddev'):
```

```
            stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
```

```
        tf.summary.scalar('stddev', stddev)
```

```
        tf.summary.scalar('max', tf.reduce_max(var))
```

```
        tf.summary.scalar('min', tf.reduce_min(var))
```

```
        tf.summary.histogram('histogram', var)
```

In [9]: *#simple scaling of inputs: limiting the data due to RAM usage issues*

```
X_train = X_train_a[:20000]/255
```

```
X_test = x_test/255
```

```
y_train = y_train_a[:20000]
```



```
y_test = y_test
```

```
In [10]: #transforming labels
```

```
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [11]: #initializing the graph
```

```
tf.reset_default_graph()
```

```
x = tf.placeholder(tf.float32, shape=(None, 32, 32, 3), name='input_x')
y = tf.placeholder(tf.float32, shape=(None, 10), name='output_y')
```

```
conv_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 3, 1], mean=0, stddev=0.1))
W = tf.Variable(tf.truncated_normal(shape=[64, num_classes], mean=0, stddev=0.1))
bias = tf.Variable(tf.truncated_normal(shape=[num_classes], mean=0, stddev=0.1))
```

```
WARNING:tensorflow:From C:\Users\User\Anaconda3\lib\site-packages\tensorflow\python\framework\
Instructions for updating:
Colocations handled automatically by placer.
```

```
In [13]: #hyperparameters
```

```
learning_rate = 1e-3
num_epochs = 30
batch_size = 128
reg_weight = 0.001
```

```
#defining the graph
```

```
conv1 = tf.nn.conv2d(x, conv_filter, strides=[1,2,2,1], padding='SAME', name="conv1")
activ1 = tf.nn.relu(conv1)
pool = tf.nn.max_pool(activ1, ksize=[1,2,2,1], strides=[1,2,2,1], padding='VALID')
flat = tf.contrib.layers.flatten(pool)
output = tf.nn.relu(tf.matmul(flat, W) + bias)
```

```
regularizer = tf.nn.l2_loss(W) + tf.nn.l2_loss(conv1)
```

```
#defintion of loss function and optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=y, logits=output))
optimizer = tf.train.RMSPropOptimizer(learning_rate=learning_rate).minimize(cost)
```

```
# Accuracy
```

```
correct_pred = tf.equal(tf.argmax(output, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')
```

```
acc_test_list = []
acc_train_list = []
acc_train_list_ = []
```

```
#suscribing tensor to tensorboard
```

```

with tf.name_scope('conv1'):
    variable_summaries(conv1)

with tf.name_scope('W'):
    variable_summaries(W)

with tf.name_scope('performance'):
    tf.summary.scalar('accuracy', accuracy)
    tf.summary.scalar('cost', cost)

#running the graph
with tf.Session() as sess:

    sess.run(tf.global_variables_initializer())

    summ_writer_train = tf.summary.FileWriter(os.path.join('summaries_simple', 'train'))
    summ_writer_test = tf.summary.FileWriter(os.path.join('summaries_simple', 'test'))

    merged = tf.summary.merge_all()

    for epoch in range(num_epochs):
        minibatches = random_mini_batches(X_train, y_train, batch_size, 1)
        list_acc_ = []

        for i, minibatch in enumerate(minibatches):
            batch_X, batch_y = minibatch

            _, cost_, acc_ = sess.run([optimizer, cost, accuracy],
                                     feed_dict = {x: batch_X,
                                                  y: batch_y})

            list_acc_.append(acc_)

        summ_test, acc_test = sess.run([merged, accuracy], feed_dict = {x: X_test,
                                                                      y: y_test})
        summ_train, acc_train = sess.run([merged, accuracy], feed_dict = {x: X_train[
                                                                      y: y_train[:10000,:]])

        acc_train_list.append(acc_train)
        acc_test_list.append(acc_test)
        acc_train_list_.append(np.mean(list_acc_))

        print("Accuracy train:", acc_train)
        print("Accuracy test:", acc_test)

        summ_writer_train.add_summary(summ_train, epoch)
        summ_writer_test.add_summary(summ_test, epoch)
        summ_writer_train.flush()
        summ_writer_test.flush()

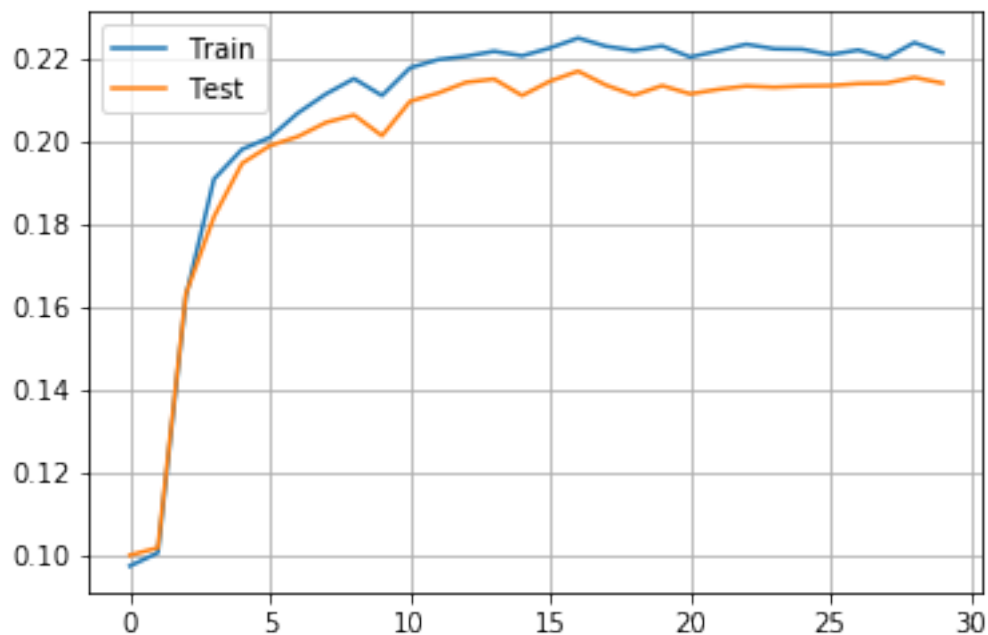
```

Accuracy train: 0.0974  
Accuracy test: 0.1  
Accuracy train: 0.1006  
Accuracy test: 0.1018  
Accuracy train: 0.1626  
Accuracy test: 0.1632  
Accuracy train: 0.1909  
Accuracy test: 0.1818  
Accuracy train: 0.1981  
Accuracy test: 0.1947  
Accuracy train: 0.201  
Accuracy test: 0.199  
Accuracy train: 0.2069  
Accuracy test: 0.2012  
Accuracy train: 0.2115  
Accuracy test: 0.2046  
Accuracy train: 0.2152  
Accuracy test: 0.2064  
Accuracy train: 0.2111  
Accuracy test: 0.2014  
Accuracy train: 0.2178  
Accuracy test: 0.2097  
Accuracy train: 0.2198  
Accuracy test: 0.2117  
Accuracy train: 0.2206  
Accuracy test: 0.2143  
Accuracy train: 0.2218  
Accuracy test: 0.2151  
Accuracy train: 0.2207  
Accuracy test: 0.2111  
Accuracy train: 0.2226  
Accuracy test: 0.2146  
Accuracy train: 0.225  
Accuracy test: 0.217  
Accuracy train: 0.223  
Accuracy test: 0.2136  
Accuracy train: 0.222  
Accuracy test: 0.2112  
Accuracy train: 0.2231  
Accuracy test: 0.2135  
Accuracy train: 0.2204  
Accuracy test: 0.2115  
Accuracy train: 0.2219  
Accuracy test: 0.2126  
Accuracy train: 0.2235  
Accuracy test: 0.2134  
Accuracy train: 0.2224  
Accuracy test: 0.2131

```
Accuracy train: 0.2223
Accuracy test: 0.2134
Accuracy train: 0.221
Accuracy test: 0.2135
Accuracy train: 0.2221
Accuracy test: 0.214
Accuracy train: 0.2202
Accuracy test: 0.2141
Accuracy train: 0.2239
Accuracy test: 0.2155
Accuracy train: 0.2215
Accuracy test: 0.2141
```

```
In [15]: plt.plot(acc_train_list)
         plt.plot(acc_test_list)
         plt.grid()
         plt.legend(("Train", "Test"))
```

```
Out[15]: <matplotlib.legend.Legend at 0x2361a8a5240>
```



### 2.1.3 More complex network

Now we run a more complex network (with two convolutional layers and two fully connected layers). We also change the activation function: instead of using ReLU, we use SeLU. We see that we achieved better results, however, the network can be further trained (with the full training

set). We limit the training set due to RAM memory usage issues in the personal computer and in laboratory.

In [21]: *#defining the graph*

```
tf.reset_default_graph()
```

```
#placeholders
```

```
x = tf.placeholder(tf.float32, shape=(None, 32, 32, 3), name='input_x')
```

```
y = tf.placeholder(tf.float32, shape=(None, 10), name='output_y')
```

```
#variables
```

```
conv1_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 3, 5], mean=0, stddev=0.1))
```

```
conv2_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 5, 10], mean=0, stddev=0.1))
```

```
W1 = tf.Variable(tf.truncated_normal(shape=[640, 100], mean=0, stddev=0.1))
```

```
bias1 = tf.Variable(tf.truncated_normal(shape=[1, 100], mean=0, stddev=0.1))
```

```
W2 = tf.Variable(tf.truncated_normal(shape=[100, num_classes], mean=0, stddev=0.1))
```

```
bias2 = tf.Variable(tf.truncated_normal(shape=[num_classes], mean=0, stddev=0.1))
```

```
#first convolutional layer
```

```
conv1 = tf.nn.conv2d(x, conv1_filter, strides=[1,1,1,1], padding='SAME', name="conv1")
```

```
conv1_bn = tf.layers.batch_normalization(conv1)
```

```
conv1_a = tf.nn.relu(conv1_bn)
```

```
conv1_pool = tf.nn.max_pool(conv1_a, ksize=[1,2,2,1], strides=[1,2,2,1], padding='VALID')
```

```
#second convolutional layer
```

```
conv2 = tf.nn.conv2d(conv1_pool, conv2_filter, strides=[1,1,1,1], padding='SAME', name="conv2")
```

```
conv2_bn = tf.layers.batch_normalization(conv2)
```

```
conv2_a = tf.nn.relu(conv2_bn)
```

```
conv2_pool = tf.nn.max_pool(conv2_a, ksize=[1,2,2,1], strides=[1,2,2,1], padding='VALID')
```

```
#fully connected layers
```

```
flat = tf.contrib.layers.flatten(conv2_pool)
```

```
fc1 = tf.matmul(flat, W1) + bias1
```

```
fc1_bn = tf.layers.batch_normalization(fc1)
```

```
fc1_a = tf.nn.relu(fc1_bn)
```

```
output = tf.matmul(fc1_a, W2) + bias2
```

In [22]: learning\_rate = 1e-3

```
num_epochs = 40
```

```
batch_size = 128
```

```
reg_weight = 0.0001
```

```
#defining cost function and optimizer
```

```
regularizer = tf.nn.l2_loss(W1) + tf.nn.l2_loss(W2) + tf.nn.l2_loss(conv1)
```

```

+ tf.nn.l2_loss(conv2)
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels= y,
                                                                logits=output))
+ reg_weight*regularizer
optimizer = tf.train.RMSPropOptimizer(learning_rate=learning_rate).minimize(cost)

# Accuracy
correct_pred = tf.equal(tf.argmax(output, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')

#susbscribing tensors to tensorboard
with tf.name_scope('conv1'):
    variable_summaries(conv1)

with tf.name_scope('conv2'):
    variable_summaries(conv2)

with tf.name_scope('W1'):
    variable_summaries(W1)

with tf.name_scope('W2'):
    variable_summaries(W2)

with tf.name_scope('performance'):
    tf.summary.scalar('accuracy', accuracy)
    tf.summary.scalar('cost', cost)

acc_test_list = []
acc_train_list = []
acc_train_list_ = []

#running the graph
with tf.Session() as sess:

    #initialinzg graph
    sess.run(tf.global_variables_initializer())

    summ_writer_train = tf.summary.FileWriter(os.path.join('summaries_complex', 'train'))
    summ_writer_test = tf.summary.FileWriter(os.path.join('summaries_complex', 'test'))

    merged = tf.summary.merge_all()

    #iterations over epochs
    for epoch in range(num_epochs):

        minibatches = random_mini_batches(X_train, y_train, batch_size, 1)
        list_acc_ = []
        for i, minibatch in enumerate(minibatches):

```

```

        batch_X, batch_y = minibatch

        _, cost_, acc_ = sess.run([optimizer, cost, accuracy],
                                   feed_dict = {x: batch_X,
                                                y: batch_y})

        list_acc_.append(acc_)

#computing the accuracy
        summ_test, acc_test = sess.run([merged, accuracy], feed_dict = {x: X_test,
                                                                           y: y_test})

        summ_train, acc_train = sess.run([merged, accuracy], feed_dict = {x: X_train[
                                                                           y: y_train[:10000,:]])

        acc_train_list.append(acc_train)
        acc_test_list.append(acc_test)
        acc_train_list_.append(np.mean(list_acc_))

#printing the accuracy
        print("Epoch ", epoch)
        print("Accuracy train:", acc_train)
        print("Accuracy test:", acc_test)

        summ_writer_train.add_summary(summ_train, epoch)
        summ_writer_test.add_summary(summ_test, epoch)
        summ_writer_train.flush()
        summ_writer_test.flush()

```

```

Epoch 0
Accuracy train: 0.3135
Accuracy test: 0.3164
Epoch 1
Accuracy train: 0.4092
Accuracy test: 0.4006
Epoch 2
Accuracy train: 0.4446
Accuracy test: 0.4329
Epoch 3
Accuracy train: 0.4476
Accuracy test: 0.4365
Epoch 4
Accuracy train: 0.4687
Accuracy test: 0.4506
Epoch 5
Accuracy train: 0.4756
Accuracy test: 0.4559
Epoch 6
Accuracy train: 0.488
Accuracy test: 0.4595

```

Epoch 7  
Accuracy train: 0.5159  
Accuracy test: 0.4803  
Epoch 8  
Accuracy train: 0.5319  
Accuracy test: 0.4882  
Epoch 9  
Accuracy train: 0.5488  
Accuracy test: 0.4988  
Epoch 10  
Accuracy train: 0.5622  
Accuracy test: 0.509  
Epoch 11  
Accuracy train: 0.565  
Accuracy test: 0.5103  
Epoch 12  
Accuracy train: 0.5845  
Accuracy test: 0.5188  
Epoch 13  
Accuracy train: 0.574  
Accuracy test: 0.5109  
Epoch 14  
Accuracy train: 0.5884  
Accuracy test: 0.5186  
Epoch 15  
Accuracy train: 0.5844  
Accuracy test: 0.5135  
Epoch 16  
Accuracy train: 0.5918  
Accuracy test: 0.5184  
Epoch 17  
Accuracy train: 0.6023  
Accuracy test: 0.5252  
Epoch 18  
Accuracy train: 0.6079  
Accuracy test: 0.5338  
Epoch 19  
Accuracy train: 0.6136  
Accuracy test: 0.536  
Epoch 20  
Accuracy train: 0.6144  
Accuracy test: 0.5368  
Epoch 21  
Accuracy train: 0.6181  
Accuracy test: 0.5376  
Epoch 22  
Accuracy train: 0.6017  
Accuracy test: 0.5233




Epoch 23  
Accuracy train: 0.6246  
Accuracy test: 0.5349  
Epoch 24  
Accuracy train: 0.606  
Accuracy test: 0.5188  
Epoch 25  
Accuracy train: 0.6235  
Accuracy test: 0.5341  
Epoch 26  
Accuracy train: 0.6153  
Accuracy test: 0.5282  
Epoch 27  
Accuracy train: 0.6112  
Accuracy test: 0.5254  
Epoch 28  
Accuracy train: 0.634  
Accuracy test: 0.5368  
Epoch 29  
Accuracy train: 0.6401  
Accuracy test: 0.5399  
Epoch 30  
Accuracy train: 0.6234  
Accuracy test: 0.5247  
Epoch 31  
Accuracy train: 0.6423  
Accuracy test: 0.5417  
Epoch 32  
Accuracy train: 0.6062  
Accuracy test: 0.5218  
Epoch 33  
Accuracy train: 0.6495  
Accuracy test: 0.5436  
Epoch 34  
Accuracy train: 0.6373  
Accuracy test: 0.5338  
Epoch 35  
Accuracy train: 0.6542  
Accuracy test: 0.5453  
Epoch 36  
Accuracy train: 0.6595  
Accuracy test: 0.5491  
Epoch 37  
Accuracy train: 0.6537  
Accuracy test: 0.5441  
Epoch 38  
Accuracy train: 0.6355  
Accuracy test: 0.5353

Epoch 39  
Accuracy train: 0.6741  
Accuracy test: 0.55

```
In [23]: plt.plot(acc_train_list)
plt.plot(acc_test_list)
plt.grid()
plt.legend(("Train", "Test"))
```

Out[23]: <matplotlib.legend.Legend at 0x2361d05d4a8>



## 2.1.4 References


- [1] Code from Minibatches taken from: [https://github.com/andersy005/deep-learning-specialization-coursera/blob/master/02-Improving-Deep-Neural-Networks/week3/Programming-Assignments/tf\\_utils.py](https://github.com/andersy005/deep-learning-specialization-coursera/blob/master/02-Improving-Deep-Neural-Networks/week3/Programming-Assignments/tf_utils.py)
- [2] Tensorboard: [https://www.tensorflow.org/guide/summaries\\_and\\_tensorboard](https://www.tensorflow.org/guide/summaries_and_tensorboard)
- [3] CIFAR 10: <https://www.cs.utoronto.ca/~kriz/cifar.html>
- [4] Data augmentation in Tensorflow (Source 1): <http://androidkt.com/tensorflow-image-augmentation-using-tf-image/>
- [5] Data augmentation in Tensorflow (Source 2): <https://medium.com/ymedialabs-innovation/data-augmentation-techniques-in-cnn-using-tensorflow-371ae43d5be9>
- [6] Sergey Ioffe and Christin Szegedy: "Batch Normalization: Accelerating Deep Netowrk Training by Reducing Internal Covariance Shift" (Available in: <https://arxiv.org/pdf/1502.03167v3.pdf>).

## Some tensorboard graphs:


### 1. Weights for the fully connected layer

W1


summaries/max\_1  
tag: W1/summaries/max\_1




summaries/mean\_1  
tag: W1/summaries/mean\_1



summaries/min\_1  
tag: W1/summaries/min\_1




summaries/stddev\_1  
tag: W1/summaries/stddev\_1




## 2. Weights for the second fully connected layer

W2


summaries/max\_1  
tag: W2/summaries/max\_1




summaries/mean\_1  
tag: W2/summaries/mean\_1



summaries/min\_1  
tag: W2/summaries/min\_1




summaries/stddev\_1  
tag: W2/summaries/stddev\_1




### 3. Weights for the filters

conv2


summaries/max\_1  
tag: conv2/summaries/max\_1




summaries/mean\_1  
tag: conv2/summaries/mean\_1



summaries/min\_1  
tag: conv2/summaries/min\_1




summaries/stddev\_1  
tag: conv2/summaries/stddev\_1



## performance


accuracy

tag: performance/accuracy



cost

tag: performance/cost




## 4. Distributions for weights in the second fully layer

W2

W2/summaries/histogram

test



W2/summaries/histogram

train

